

# Data Architecture & Advanced SQL

JOIN • CTE • Window Functions • Time Travel  
Hands-on: Event Pipeline (orders, sessions, payments)

# Data Architecture & Advanced SQL



Course: Business Analytics



Topics: JOIN • CTE • Window Functions • Time Travel



Hands-on: Event Pipeline (orders, sessions, payments)

# Agenda

- 1) Data architecture for analytics: OLTP vs OLAP, Star/Snowflake, Lakehouse, Medallion
- 2) Advanced JOINs: inner/outer, cross, semi/anti, join strategies
- 3) CTE & Recursive CTE: readability, reuse, graph/tree examples
- 4) Window functions: ranking, moving average, frames, sessionization
- 5) Time Travel: query historical data (Snowflake, BigQuery, Delta Lake)
- 6) Hands-on: Event pipeline queries (orders, sessions, payments)


# Learning Outcomes

- Explain core data architecture patterns and trade-offs
- Write correct and performant JOINs (including semi/anti)
- Use CTEs and recursive CTEs to structure complex logic
- Apply window functions for rolling metrics and cohort analysis
- Query historical states using time travel features
- Build an end-to-end SQL pipeline over events data




# Part I — Data Architecture for Analytics

- OLTP vs OLAP workloads: read/write patterns
- Normalization vs denormalization
- Dimensional modeling (Star & Snowflake)
- Lakehouse & Medallion architecture
- Batch vs streaming; CDC & event-driven pipelines



## OLTP vs OLAP

- OLTP: many small writes, normalized schema, point lookups
- OLAP: fewer large reads, denormalized schema, columnar storage
- Implication: separate systems, or lakehouse with serving layers



# Dimensional Modeling: Star Schema



Fact table: transactional grain (e.g., orders, payments)




Dimension tables: customers, products, time, channel



Benefits: simplicity, predictable JOINs, good for BI


## Dimensional Modeling: Snowflake Schema

- Normalized dimensions (e.g., product → category → department)
- Pros: reduced redundancy; Cons: more joins and complexity
- Use when dimension updates are frequent and large



## SCD Type 2 (Slowly Changing Dimensions)

- Maintain history via versioned rows with validity ranges
- Columns: valid\_from, valid\_to, is\_current
- Join fact → dimension using time-sensitive key + range




# SCD Type 2 (Slowly Changing Dimensions) — Code

```
-- Example SCD2 join (PostgreSQL-ish)
SELECT f.order_id, d.customer_tier
FROM fact_orders f
JOIN dim_customer_scd d
  ON f.customer_id = d.customer_id
  AND f.order_ts >= d.valid_from
  AND (d.valid_to IS NULL OR f.order_ts < d.valid_to);
```

# Lakehouse & Medallion

- Bronze: raw ingests (append-only, minimal transforms)
- Silver: cleaned & conformed (dedupe, schema), ready for JOINs
- Gold: curated marts for BI/ML (star schemas, aggregates)



# Batch vs Streaming

- Batch: simpler orchestration, higher latency
- Streaming: lower latency, complexity (exactly-once, upserts)
- Micro-batch (e.g., Spark Structured Streaming) bridges both

# CDC (Change Data Capture) & Event-Driven

- Capture inserts/updates/deletes from OLTP (e.g., Debezium)
- Publish to topics; process into lakehouse layers
- Upsert into silver; maintain gold aggregates

# Event Data Model (orders, sessions, payments)

- Events: event\_time, event\_type, user\_id/session\_id
- orders(order\_id, user\_id, session\_id, amount, status, created\_at)
- payments(payment\_id, order\_id, amount, method, status, paid\_at)
- sessions(session\_id, user\_id, start\_at, end\_at, channel)

## Part II — Advanced JOINS

- Inner/Left/Right/Full/Cross
- Semi/Anti joins (EXISTS / NOT EXISTS)
- Join keys, NULL handling, duplicates, cardinality
- Join order & execution strategies

# Equi-Join vs Non-Equi Join

Equi-join on keys:  
common for  
dimensions

Non-equi: range  
joins (e.g., SCD2  
validity windows)

# Inner vs Outer Joins

INNER: rows that  
match on both  
sides

LEFT/RIGHT:  
preserve one side  
with NULLs for  
missing matches

FULL: preserve  
both sides

# Inner vs Outer Joins — Code

```
-- Outer join example: orders without payments (left anti via WHERE)
SELECT o.order_id, o.amount
FROM orders o
LEFT JOIN payments p ON p.order_id = o.order_id
WHERE p.order_id IS NULL;
```

# Semi & Anti Joins

SEMI: rows from A  
with at least one  
match in B  
(EXISTS)

ANTI: rows from A  
with no matches  
in B (NOT EXISTS)

Useful for integrity  
checks and funnels

# Semi & Anti Joins — Code

```
-- Semi join: orders that have payments
```

```
SELECT o.order_id
FROM orders o
WHERE EXISTS (SELECT 1 FROM payments p WHERE p.order_id = o.order_id);
```

```
-- Anti join: orders lacking payments
```

```
SELECT o.order_id
FROM orders o
WHERE NOT EXISTS (SELECT 1 FROM payments p WHERE p.order_id = o.order_id);
```

# Join Gotchas & NULLs

NULLs do not equal anything (even NULL) in equality joins

Coalesce or surrogate keys for null-safe comparisons

Beware duplicated keys → row explosion

# Join Gotchas & NULLs — Code

```
-- Null-safe equality (BigQuery)
SELECT *
FROM a
LEFT JOIN b
ON IFNULL(a.key, '__NULL__') = IFNULL(b.key, '__NULL__');
```

# Join Strategies & Performance

Broadcast/small  
table join vs  
shuffle join

Hash join vs sort-  
merge join

Pushdown filters,  
pre-aggregate,  
selective  
projections

# Part III — CTEs & Recursive CTEs

CTE: improves readability;  
can aid reuse

Recursive CTE: trees/graphs,  
hierarchies, sequences

Materialization behavior  
varies by engine (cost-based)

# CTE for Organization

Break complex logic into named steps (bronze → silver → gold)

Avoid repeated subqueries;  
annotate transformations

# CTE for Organization — Code

```
WITH clean_orders AS (
    SELECT order_id, user_id, amount, created_at::date AS order_date
    FROM raw_orders
    WHERE status = 'confirmed'
), daily_rev AS (
    SELECT order_date, SUM(amount) AS revenue
    FROM clean_orders GROUP BY 1
)
SELECT * FROM daily_rev ORDER BY order_date;
```

# Recursive CTE — Sequence & Tree

Generate series;  
walk parent-child  
relations

Add a termination  
condition to avoid  
infinite loops

# Recursive CTE — Sequence & Tree — Code

```
-- Tree traversal (PostgreSQL-style)
WITH RECURSIVE org AS (
    SELECT id, parent_id, name, 0 AS depth
    FROM dept WHERE parent_id IS NULL
    UNION ALL
    SELECT d.id, d.parent_id, d.name, o.depth + 1
    FROM dept d
    JOIN org o ON d.parent_id = o.id
)
SELECT * FROM org ORDER BY depth, name;
```

# Part IV — Window Functions

Compute metrics  
over partitions  
without collapsing  
rows

RANK /  
DENSE\_RANK /  
ROW\_NUMBER

SUM / AVG over  
frames for moving  
windows

# Window Syntax

- `function(...) OVER (PARTITION BY ... ORDER BY ... [FRAME])`
- Frames: ROWS vs RANGE vs GROUPS
- Default frame varies by engine; specify explicitly

# Window Syntax — Code

```
-- 7-day rolling revenue (date-grain)
SELECT order_date,
       SUM(amount) OVER (
           ORDER BY order_date
           ROWS BETWEEN 6 PRECEDING AND CURRENT ROW
       ) AS rev_7d
FROM daily_orders;
```

# Ranking Examples

Top-N per category via  
ROW\_NUMBER()

Percentiles with  
NTILE() or  
PERCENT\_RANK()

# Ranking Examples — Code

```
-- Top 3 products per category
SELECT *
FROM (
    SELECT category, product_id, sales,
           ROW_NUMBER() OVER (PARTITION BY category ORDER BY sales DESC) AS rn
    FROM product_sales
) s
WHERE rn <= 3;
```

# Sessionization with Windows

- Assign sessions by gaps between events
- Use LAG() to detect gap > threshold
- Some engines have SESSION windows; generic SQL shown

# Sessionization with Windows — Code

```
-- Sessionize: start new session if gap > 30 minutes
WITH g AS (
    SELECT user_id, event_ts,
           CASE WHEN event_ts - LAG(event_ts) OVER (PARTITION BY user_id ORDER BY
event_ts)
                 > INTERVAL '30 minutes' OR LAG(event_ts) IS NULL
                 THEN 1 ELSE 0 END AS is_new
      FROM events
),
s AS (
    SELECT user_id, event_ts,
           SUM(is_new) OVER (PARTITION BY user_id ORDER BY event_ts
                           ROWS UNBOUNDED PRECEDING) AS session_num
      FROM g
)
SELECT user_id, MIN(event_ts) AS session_start, MAX(event_ts) AS session_end,
       COUNT(*) AS events
  FROM s
 GROUP BY user_id, session_num;
```

# Cohort & Retention with Windows

- Cohort by first order month; retention by return months
- Use MIN() OVER for cohort; use DATEDIFF for offsets

# Cohort & Retention with Windows — Code

```
WITH firsts AS (
    SELECT user_id,
        DATE_TRUNC('month', MIN(order_ts)) OVER (PARTITION BY user_id) AS cohort_month
        order_ts
    FROM orders
),
ret AS (
    SELECT user_id, cohort_month,
        DATE_DIFF('month', cohort_month, DATE_TRUNC('month', order_ts)) AS month_idx
    FROM firsts
)
SELECT cohort_month, month_idx, COUNT(DISTINCT user_id) AS users
FROM ret
GROUP BY 1,2 ORDER BY 1,2;
```

# Part V — Time Travel

Query historical table  
versions to audit or recover

Common in Snowflake,  
BigQuery, Delta Lake

Limits & retention vary by  
platform

# Time Travel — Snowflake

- Point-in-time query with AT/BEFORE clause
- Retention: typically 1–90 days (account setting)

# Time Travel — Snowflake — Code

```
-- Snowflake
SELECT * FROM sales AT (TIMESTAMP => TO_TIMESTAMP('2025-08-31 00:00:00'));
-- or by offset seconds
SELECT * FROM sales AT (OFFSET => -3600);
```

# Time Travel — BigQuery

- Use `FOR SYSTEM_TIME AS OF timestamp`
- Default retention ~7 days for table snapshots

# Time Travel — BigQuery — Code

```
-- BigQuery
SELECT * FROM `proj.ds.sales` FOR SYSTEM_TIME AS OF
TIMESTAMP_SUB(CURRENT_TIMESTAMP(), INTERVAL 1 DAY);
```

# Time Travel — Delta Lake

Query by VERSION  
AS OF or  
TIMESTAMP AS OF

Inspect history via  
DESCRIBE  
HISTORY

# Time Travel — Delta Lake — Code

```
-- Delta Lake
SELECT * FROM delta.`/tables/sales` VERSION AS OF 12;
-- Or:
SELECT * FROM delta.`/tables/sales` TIMESTAMP AS OF '2025-08-31T00:00:00Z';
```

# Part VI — Hands-on: Event Pipeline

Dataset: orders, payments, sessions

Tasks: funnel, anomalies, rolling metrics, time-travel checks

Environment-agnostic SQL  
(PostgreSQL/SparkSQL/BigQuery-style)

# Create Bronze (raw → cleaned)

- Deduplicate, cast types, basic validation

# Create Bronze (raw → cleaned) – Code

```
WITH o AS (
    SELECT DISTINCT order_id, user_id, session_id, amount::NUMERIC, status, created_at::timestamp AS order_ts
    FROM raw_orders
),
p AS (
    SELECT DISTINCT payment_id, order_id, amount::NUMERIC, status, paid_at::timestamp AS paid_ts
    FROM raw_payments
),
s AS (
    SELECT DISTINCT session_id, user_id, start_at::timestamp AS session_start, end_at::timestamp AS session_end, channel
    FROM raw_sessions
)
SELECT COUNT(*) FROM o; -- sanity checks
```

# Silver: Conform & Joinable Tables

- Normalize status values; ensure referential integrity

# Silver: Conform & Joinable Tables – Code

```
WITH clean_orders AS (
    SELECT *, LOWER(status) AS status_norm
    FROM o
    WHERE amount >= 0 AND order_ts IS NOT NULL
), clean_payments AS (
    SELECT *, LOWER(status) AS status_norm FROM p
), clean_sessions AS (
    SELECT * FROM s WHERE session_start <= COALESCE(session_end, session_start + INTERVAL '8 hours')
)
SELECT COUNT(*) FROM clean_orders;
```

# Gold: Core Marts

- fact\_orders, fact\_payments, dim\_user, dim\_channel
- Daily revenue, conversion rates, AOV

# Gold: Core Marts — Code

```
-- Daily revenue & orders
WITH daily AS (
    SELECT DATE_TRUNC('day', order_ts) AS d, SUM(amount) AS revenue, COUNT(*) AS orders
    FROM clean_orders WHERE status_norm='confirmed' GROUP BY 1
)
SELECT d, revenue, orders, revenue::NUMERIC / NULLIF(orders,0) AS aov
FROM daily ORDER BY d;
```

# Funnel: session → order → payment

- Identify drop-offs with left/semi/anti joins

# Funnel: session → order → payment — Code

```
WITH session_to_order AS (
    SELECT s.session_id, COUNT(DISTINCT o.order_id) AS orders
    FROM clean_sessions s
    LEFT JOIN clean_orders o ON o.session_id = s.session_id
    GROUP BY 1
),
order_to_payment AS (
    SELECT o.order_id, EXISTS(SELECT 1 FROM clean_payments p
    WHERE p.order_id=o.order_id AND p.status_norm='succeeded') AS paid
    FROM clean_orders o
)
SELECT
    (SELECT COUNT(*) FROM clean_sessions) AS sessions,
    (SELECT COUNT(DISTINCT order_id) FROM clean_orders) AS orders,
    (SELECT COUNT(*) FROM order_to_payment WHERE paid) AS paid_orders;
```

# Payment Mismatch Checks

- Amount mismatches, multiple payments per order, failed payments

# Payment Mismatch Checks — Code

```
-- Mismatch report
SELECT o.order_id, o.amount AS order_amount, SUM(p.amount) AS paid_amount
FROM clean_orders o
LEFT JOIN clean_payments p ON p.order_id=o.order_id AND p.status_norm='succeeded'
GROUP BY 1, o.amount
HAVING SUM(p.amount) IS DISTINCT FROM o.amount;
```

# Rolling KPIs with Windows

- 7-day revenue, 30-day conversion rate

# Rolling KPIs with Windows — Code

```
-- 7-day rolling revenue & orders
WITH daily AS (
    SELECT DATE_TRUNC('day', order_ts) AS d, SUM(amount) AS revenue, COUNT(*) AS orders
    FROM clean_orders WHERE status_norm='confirmed' GROUP BY 1
)
SELECT d,
    SUM(revenue) OVER (ORDER BY d ROWS BETWEEN 6 PRECEDING AND CURRENT ROW) AS rev_7d,
    SUM(orders) OVER (ORDER BY d ROWS BETWEEN 6 PRECEDING AND CURRENT ROW) AS orders_7d
FROM daily ORDER BY d;
```

# User Cohorts & LTV

- First-order cohort, months since cohort, cumulative revenue

# User Cohorts & LTV — Code

```
WITH first_order AS (
    SELECT user_id, MIN(order_ts) AS first_ts
    FROM clean_orders GROUP BY 1
),
cohort_orders AS (
    SELECT o.user_id,
        DATE_TRUNC('month', f.first_ts) AS cohort_month,
        DATE_DIFF('month', DATE_TRUNC('month', f.first_ts),
DATE_TRUNC('month', o.order_ts)) AS m,
        o.amount
    FROM clean_orders o JOIN first_order f USING(user_id)
)
SELECT cohort_month, m, SUM(amount) AS revenue
FROM cohort_orders
GROUP BY 1,2 ORDER BY 1,2;
```

# Anomaly Detection (SQL-first)

- Z-score vs global mean; quick-and-dirty monitoring

# Anomaly Detection (SQL-first) — Code

```
WITH daily AS (
    SELECT DATE_TRUNC('day', order_ts) AS d, SUM(amount) AS revenue
    FROM clean_orders WHERE status_norm='confirmed' GROUP BY 1
),
stats AS (
    SELECT AVG(revenue) AS mu, STDDEV_SAMP(revenue) AS sigma FROM daily
)
SELECT d, revenue, (revenue - mu)/NULLIF(sigma,0) AS z
FROM daily CROSS JOIN stats
ORDER BY d;
```

# Time Travel in the Pipeline

- Audit historical states of facts/dimensions
- Compare to current to detect retroactive changes

# Time Travel in the Pipeline — Code

```
-- Snowflake: yesterday's snapshot vs today
WITH y AS (
    SELECT * FROM fact_orders AT (OFFSET => -86400)
), t AS (
    SELECT * FROM fact_orders
)
SELECT COALESCE(y.order_id, t.order_id) AS order_id,
       y.amount AS amount_y, t.amount AS amount_t
FROM y FULL JOIN t USING(order_id)
WHERE y.amount IS DISTINCT FROM t.amount;
```

# End-to-End Example Query

- From bronze to gold in one CTE pipeline

# End-to-End Example Query — Code

```
WITH o AS (
    SELECT DISTINCT order_id, user_id, session_id, amount:::NUMERIC,
    LOWER(status) AS st, created_at:::timestamp AS ts FROM raw_orders
),
p AS (
    SELECT DISTINCT order_id, amount:::NUMERIC, LOWER(status)
    AS st, paid_at:::timestamp AS ts FROM raw_payments
),
s AS (
    SELECT DISTINCT session_id, user_id, start_at:::timestamp AS start_ts
    FROM raw_sessions
),
confirmed AS (SELECT * FROM o WHERE st='confirmed'),
paid AS (SELECT order_id, SUM(amount) AS paid_amount FROM p
WHERE st='succeeded' GROUP BY 1),
joined AS (
    SELECT c.order_id, c.user_id, c.session_id, c.amount, c.ts
    AS order_ts, COALESCE(p.paid_amount,0) AS paid_amount FROM confirmed c
    LEFT JOIN paid p USING(order_id)
)
SELECT *,
    CASE WHEN paid_amount >= amount THEN 'paid' ELSE 'unpaid' END AS payment_status
FROM joined;
```

# Exercises (1/2)

- Compute daily conversion rate: sessions → orders → paid
- Find orders with payment attempts > 1 and flag anomalies
- Top-3 channels by 7-day rolling revenue

## Exercises (2/2)

- Build a cohort table by signup month and compute 90-day LTV
- Use time travel to compare today's gold mart vs 7 days ago
- Recursive CTE: expand a product category tree and roll up sales

# Summary & What's Next

- We covered data architecture patterns for analytics
- Advanced SQL: JOINs, CTEs, windows, time travel
- Hands-on pipeline for events data
- Next: KPI Tree, Experimentation, and Causal Thinking

# References

- Kimball & Ross — The Data Warehouse Toolkit
- Snowflake, BigQuery, Delta Lake official docs  
(time travel, windows)
- Spark SQL, PostgreSQL docs (CTE, window functions)

## Extra Example — Window Frame Deep Dive

- Compare ROWS vs RANGE on date buckets
- Demonstrate GROUPS frame for ties (if supported)

# Extra Example — Window Frame Deep Dive — Code

```
-- RANGE vs ROWS example (beware duplicates on ORDER BY key)
SELECT d, x,
       SUM(x) OVER (ORDER BY d ROWS BETWEEN 2 PRECEDING AND CURRENT ROW) AS rows_win,
       SUM(x) OVER (ORDER BY d RANGE BETWEEN INTERVAL '2 day'
PRECEDING AND CURRENT ROW) AS range_win
FROM series;
```

