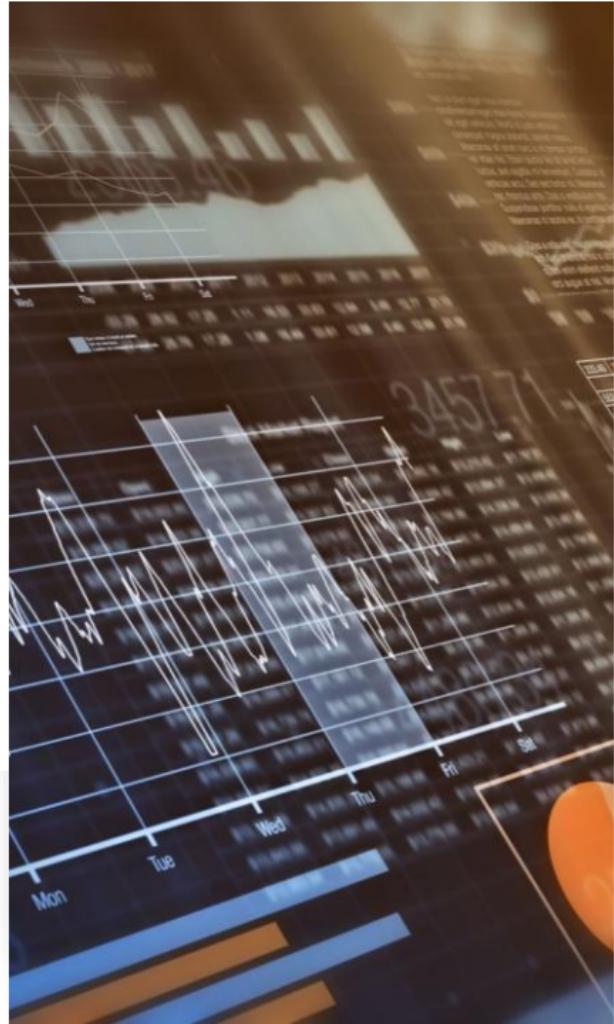


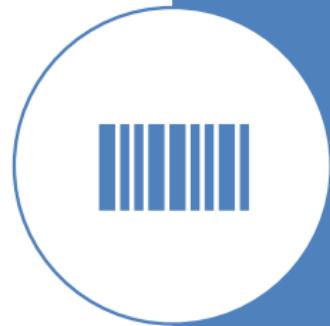
# Business Data Analytics – Lecture 3: Data Quality & Feature Engineering

Assoc. Prof. Nguyen Binh Minh, Ph.D.



# Learning Objectives (CLO)

- Understand data quality dimensions and measurement.
- Diagnose and handle missingness.
- Detect and avoid data leakage.
- Encode categorical variables and scale numeric features.
- Engineer features and perform feature selection.
- Build reproducible, leakage-safe pipelines.



# Agenda & Timing

- 1) Data quality (40')
- 2) Missing data (35')
- 3) Data leakage (25')
- 4) Categorical encoding (35')
- 5) Scaling & outliers (25')
- 6) Feature engineering & selection (40')
- 7) Pipeline & mini-case (30')
- 8) Short quiz & wrap-up (10')

# Running Case: Telco Churn Prediction



Dataset: 10k customers, 35 variables (demographics, plan, usage, label).



Goal: predict churn in the next 30 days.



Challenges: missing income/tenure, leakage (post-event features),



High cardinality `plan_id`, mixed scales.

## What is Data Quality?



Accuracy,  
Completeness,  
Consistency.



Validity, Timeliness,  
Provenance.



Direct impact on  
model performance  
and ops cost.

# Data Quality Scorecard

Example KPIs: % missing < 2%/feature;  
duplicate rate < 0.5%; out-of-range < 0.1%.

Data SLA: daily H+1 before 08:00; stable  
schema outside maintenance windows.

# Data Quality Lifecycle

Ingest → Profile →  
Clean → Validate →  
Monitor → Feedback.

Combine rule-based  
checks with  
descriptive statistics.

# Quick Data Profiling

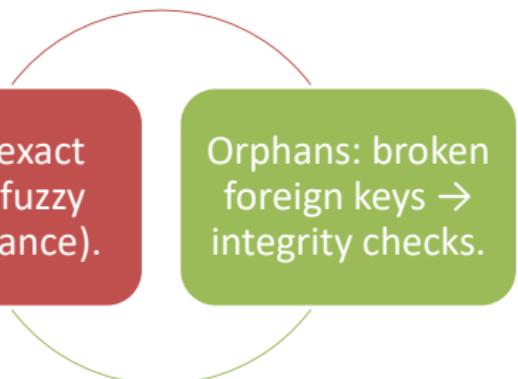
Distributions,  
min/max, outliers,  
cardinality, missing  
map.

Visuals: histogram,  
boxplot, correlation  
heatmap, category  
bars.

# Duplicates & Orphan Records

Dedup: exact  
keys vs. fuzzy  
(edit distance).

Orphans: broken  
foreign keys →  
integrity checks.

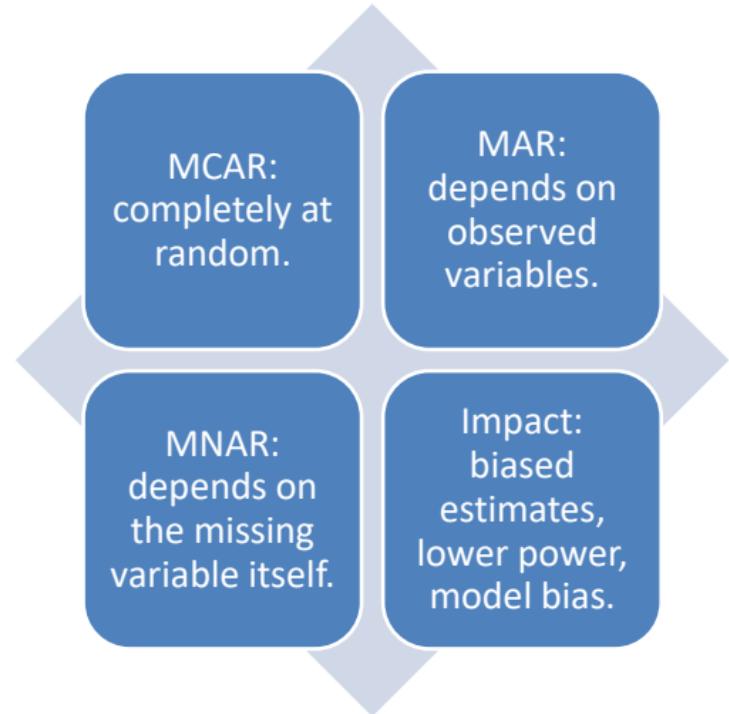


## Units & Value Domains

Normalize currencies,  
percentages, and  
units ( $\text{km} \leftrightarrow \text{m}$ ).

Validity rules (e.g.,  
 $\text{age} \in [0, 120]$ ).

# Missingness Mechanisms



## Diagnosing Missingness

Column/row missing maps; relation between missingness and label.

Little's MCAR idea; compare distributions in missing vs. non-missing groups.

## Example: Missing by Segment

Income missing more in Prepaid;  
Tenure missing in new customers.

Likely MAR by plan\_type; possible  
MNAR for income.

# Missingness Strategies — Overview

Row/column drop if excessive missingness or low-value feature.

Imputation:  
mean/median/mode,  
group-wise, kNN, MICE,  
model-based, business  
rules.

Add an `is_missing` indicator.

# Simple & Group-wise Imputation

Median/Mode  
robust to  
outliers.

Group-wise  
medians by  
`plan_type`.

# kNN Imputation

- Find k nearest neighbors → average/majority.
- Pros: preserves multivariate structure; Cons: compute-heavy, scale-sensitive.
- Apply scaling inside a pipeline to avoid leakage.

# MICE (Multiple Imputation by Chained Equations)

- Model each incomplete feature iteratively until convergence.
- Produce multiple datasets and pool estimates.
- Best for MAR; preserves correlation structure.

# Business- and Time-aware Imputation

- Time series: forward/backward fill, interpolation, seasonal Kalman.
- Domain rules:  $\text{income} = \text{base salary} \times \text{factor}$ ; apply caps.

# Assessing Imputation Impact

- Compare model metrics (AUC/RMSE) across strategies.
- Check distribution drift after imputation; keep predictive signal.

# Common Pitfalls in Missingness

- Global-mean imputation using full data → leakage.
- Forgetting the `is_missing` flag.
- Imputing before train/test split → leakage!

# Data Leakage — Definition & Impact

- Information from future/test creeps into training.
- Over-optimistic training scores, sharp drop in production.

# Leakage Types

- Target leakage: features contain the label or its consequences.
- Train–test contamination: preprocessing fitted on full data.
- Temporal leakage: using post-hoc features for forecasting.

# Example of Target Leakage (Churn)

- Feature 'refund\_in\_30d' appears in training but occurs after churn.
- Mitigation: drop or time-shift to before the prediction point.

# Avoiding Temporal Leakage

- Time-based splits (TimeSeriesSplit/rolling origin).
- Freeze feature windows to pre-T0 data only.

# Anti-Leakage Pipeline Principles

- Fit preprocessing only on training folds.
- Use Pipeline/ColumnTransformer.
- Perform cross-validation on the full pipeline.

# Data Splitting Strategies

- Holdout, KFold, StratifiedKFold.
- TimeSeriesSplit for temporal data.
- Group-aware split when multiple rows per customer.

# Categorical Encoding — Overview

- One-Hot, Ordinal, Target/Mean, WoE, Hashing, Frequency/Count.
- Issues: high cardinality, rare categories, target-encoding leakage.

# One-Hot Encoding (OHE)

- Pros: simple, no order assumption.
- Cons: dimensionality explosion with high cardinality.
- Tips: group rare levels to 'Other'; use hashing if needed.



# Ordinal Encoding

Map categories to integers with true ordering ( $S < M < L < XL$ ).

Avoid for unordered categories.

# Target/Mean Encoding (with CV)

- Replace category with target mean (with smoothing).
  - Compute inside CV loops to prevent leakage.
  - Add noise/smoothing to reduce overfit.
-

# Weight of Evidence (WoE)

- $\text{WoE} = \ln((\text{Good}_i/\text{Bad}_i)/(\text{Good}/\text{Bad}))$  — widely used in credit risk.
- Pros: monotonicity, interpretability; Cons: requires binning and enough samples.

# Hashing Trick

- Map categories to a fixed space via hashing.
- Pros: handles high cardinality and streaming; Cons: collisions, less interpretable.

# High-Cardinality Tactics

Collapse rare levels;  
frequency/count encoding.

CV-safe target encoding; or hashing + embeddings.

# Text Features (quick)

Bag-of-Words/TF-IDF; n-grams;  
remove stopwords.

Pair with linear/logistic models or tree-based learners.

# Entity Embeddings (idea)

Learn dense representations of categories with neural nets.

Great for high-cardinality interactions.

# Why & When to Scale

Scale-sensitive  
algorithms: kNN, k-  
means, SVM, PCA,  
L1/L2 regression.

Tree-based models less  
sensitive but scaling  
helps when mixing  
models.

## Scaling Techniques

---

StandardScaler:  $(x - \mu) / \sigma$ .

---

MinMaxScaler: [0,1].

---

RobustScaler: median/IQR  
for outlier robustness.

---

Always fit on train,  
transform train/test.

## Power / Variance- Stabilizing Transforms

---

Log, Box-Cox ( $x > 0$ ),  
Yeo-Johnson  
( $x \in \mathbb{R}$ ).

---

Make distributions  
closer to normal;  
reduce skew.

Quantile /  
Rank-  
Gauss  
Transform

---

Map to  
uniform/normal  
distributions.

---

Useful for linear  
models and  
normality  
assumptions.

# Outliers — Detection

---

Z-scores, IQR  
( $Q1 - 1.5 \times IQR$ ,  
 $Q3 + 1.5 \times IQR$ ).

---

Model-based: Isolation  
Forest, LOF.

---

Note: outliers may carry  
business signal.

## Outliers

---

### — Handling

---

Capping/Winsorization;  
log transforms; robust  
loss.

---

Conditional removal  
with domain  
investigation.

## Feature Transformations Toolkit

---

Binning (equal-width/frequency), interactions, polynomials.

---

Ratios/Rates (e.g., ARPU = revenue/usage).

---

Time-normalized rates.

# Date– Time Features

Calendar:  
day/week/month/quarter/year;  
holiday flags.

Seasonality with sine–cosine for  
cycles (hour/day).

Recency/freshness: days since  
last activity.

Geo  
Features (if  
coordinates)

---

Haversine distances,  
regional clustering,  
point density.

---

Grid encoding  
(geohash) for  
regional aggregation.

## Aggregation / Group-by Features

---

Counts/frequencies per user/plan.

---

Group means/medians/std; shares/ratios.

---

Leave-one-out strategies to reduce leakage.

## Time- Series Features

---

Lags ( $t-1$ ,  $t-7$ ,  $t-30$ ), rolling mean/std/max/min.

---

Exponential moving features;  
differencing.

Text/NLP  
Features  
(extended)

---

Sentiment scores,  
keyword flags  
(complaints).

---

TF-IDF + SVD (LSA)  
for dimension  
reduction.

# Feature Selection

---

## — Overview

---

Filter: Chi-square,  
ANOVA, Mutual  
Information.

---

Wrapper: RFE, SFS/SBS.

---

Embedded: L1/L2, tree-  
based importance, SHAP.

# Filter Methods

---

Chi-square for categorical vs. target.

---

ANOVA F-test for numeric vs. class.

---

Mutual Information for non-linear dependencies.

# Wrapper Methods

RFE with estimator;  
choose k via CV.

Trade-off:  
accuracy vs.  
compute time.

## Embedded & Regularization

---

Lasso (L1) for sparsity;  
Ridge (L2) for variance reduction.

---

Elastic Net balances L1/L2.

---

Tree-based importance:  
Gini/Permutation.

# Dimensionality Reduction

PCA preserves variance in orthogonal components.

t-SNE/UMAP for visualization; avoid direct forecasting use.

# Golden Rules for Pipelines

---

Fit

Never fit preprocessing on the full dataset.

Use

Use ColumnTransformer for per-type pipelines.

Wrap

Wrap hyperparameter search around the pipeline.

# Pipeline Diagram (concept)

- Raw → Split → (Impute → Encode → Scale → FeatureSelect) → Model → Eval.
- For temporal data: enforce pre-T0 feature windows.

# Sklearn Pipeline (leakage-safe) — Pseudocode

```
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder, RobustScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import StratifiedKFold, cross_val_score
from category_encoders.target_encoder import TargetEncoder

num_pipe = Pipeline([
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", RobustScaler())
])

cat_high_pipe = Pipeline([
    ("target_enc", TargetEncoder())
])

cat_low_pipe = Pipeline([
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

pre = ColumnTransformer([
    ("num", num_pipe, num_cols),
    ("cat_high", cat_high_pipe, high_card_cols),
    ("cat_low", cat_low_pipe, low_card_cols)
])

clf = Pipeline([
    ("preprocess", pre),
    ("model", LogisticRegression(max_iter=200))
])

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
score = cross_val_score(clf, X, y, cv=cv, scoring="roc_auc")
```

# Model Evaluation after Preprocessing

- Metrics: ROC-AUC/PR-AUC (class imbalance), Brier score, KS.
- Compare impute/encode/scale strategies via consistent CV.

# Data Validation & Monitoring

- Schema contracts (types, ranges, constraints).
- Pre-deploy checks: null %, ranges, cardinality, drift.
- Production monitoring with alerts and thresholds.

# Data Contracts & Schema Evolution

- Agreements between data producers and consumers: fields, definitions, SLAs.
- Schema versioning and backfill strategies.

# Reproducibility & Experiment Tracking

- Random seeds, environment files.
- Track params, metrics, artifacts, and data snapshots.

# Common Anti-patterns

- Future-derived features; preprocessing on full data.
- Unbounded one-hot; ignoring rare categories.
- No `is_missing` flags; ignoring group-aware splits.

# Mini-case: Churn — Part 1 (Profiling)

- Task: profile and define a data quality checklist.
- Deliverable: DQ metrics table for top 10 columns.

# Mini-case: Churn — Part 2 (Missing & Encoding)

- Choose strategies for income and tenure imputation.
- Encode plan\_id (high cardinality) with CV-safe target encoding.

# Mini-case: Churn — Part 3 (Scaling & Model)

- Compare Logistic (scaled) vs. XGBoost (no scaling needed).
- Report ROC-AUC/PR-AUC and calibration plot.

# In-class Hands-on

- 1) Build the leakage-safe pipeline as shown.
- 2) Run 5-fold CV and log results.
- 3) Write five lines summarizing trade-offs.

# Summary — Key Takeaways

Data quality sets the ceiling for model performance.

Leakage-safe pipelines are mandatory.

Encoding/Scaling must match the model context.

Feature engineering blends domain knowledge and controlled experiments.

# Appendix

---

## —

## Formulas & Notes

---

Standardize:  $z = (x - \mu) / \sigma.$

---

IQR =  $Q3 - Q1$ ; capping at  $[Q1 - 1.5 * IQR, Q3 + 1.5 * IQR].$

---

WoE/IV basics; time cyclic encoding:  
 $\sin(2\pi t/T), \cos(2\pi t/T).$

# Appendix

## — Group- wise

## Imputation (pandas)

- `X['income'] = X.groupby('plan_type')['income'].transform(lambda s: s.fillna(s.median()))`
- `X['income_is_missing'] = X['income'].isna().astype(int)`

# Appendix — CV-safe Target Encoding

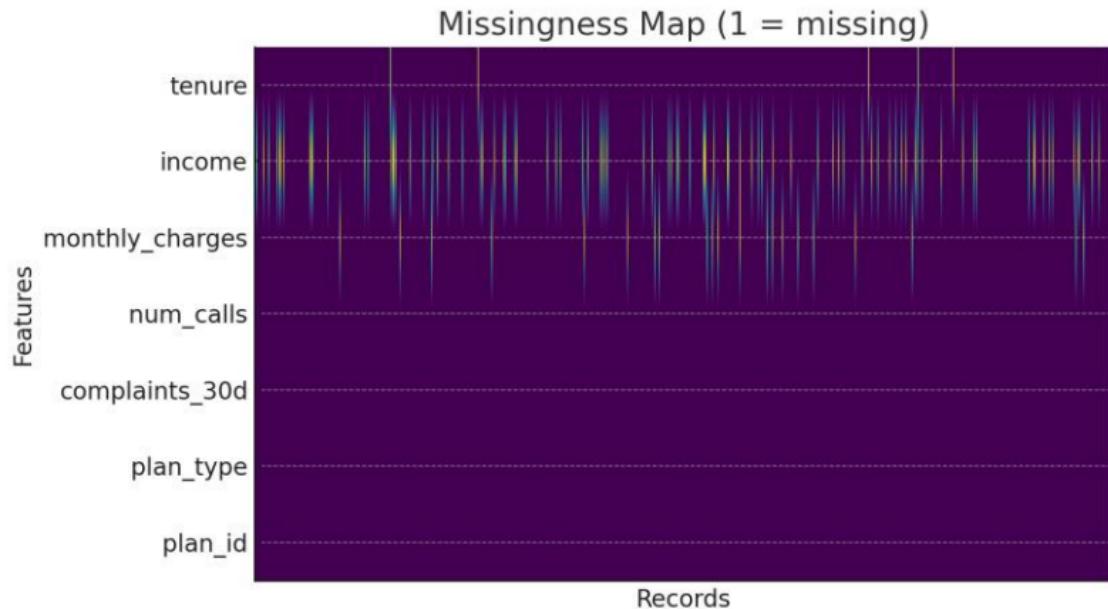
```
from category_encoders.target_encoder import  
TargetEncoder  
  
from sklearn.pipeline import Pipeline  
from sklearn.compose import ColumnTransformer  
from sklearn.linear_model import LogisticRegression  
  
te = TargetEncoder(cols=['plan_id'], smoothing=0.2)  
model = Pipeline([  
    ('pre', ColumnTransformer([('plan', te, ['plan_id'])],  
    remainder='drop')),  
    ('clf', LogisticRegression(max_iter=200))  
])  
# Use within cross_val_score/cross_val_predict to  
# avoid leakage
```

# Recommended References

- Kuhn & Johnson (2019) — Feature Engineering and Selection.
- Pang-Ning Tan et al. — Introduction to Data Mining.
- scikit-learn docs: Pipeline, ColumnTransformer, Imputation, Encoding, Scaling.
- Credit scoring resources: monotone binning, WoE/IV.

# Visual Example – Missingness Map

Synthetic telco-like dataset (1 = missing)



# Visual Example — Outliers

Boxplot of Monthly Charges

